


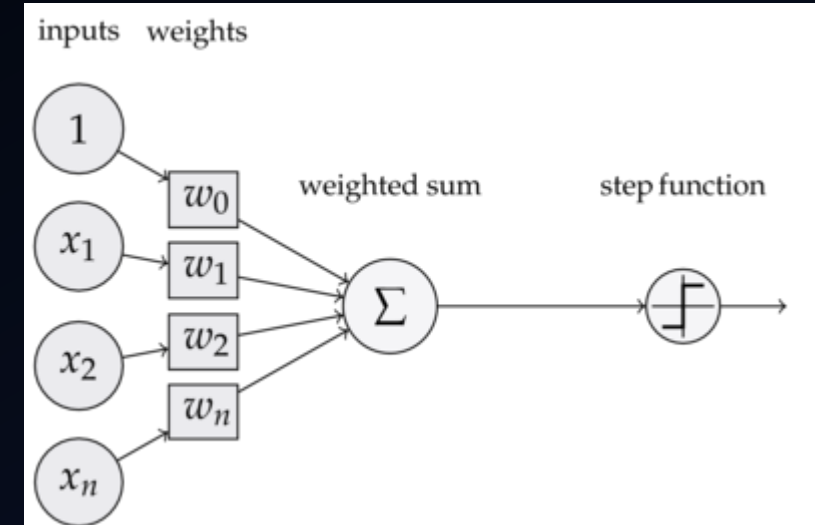
# Neural Networks and Deep Learning

Joshua Achiam

# Outline

## 1. Fundamentals


1. Perceptrons
2. Early Neural Networks
3. Backpropagation and SGD
4. Universal Approximation Theorem
5. Convolutional Networks
6. Recurrent Networks



# Outline

## 2. Recent History of Deep Learning

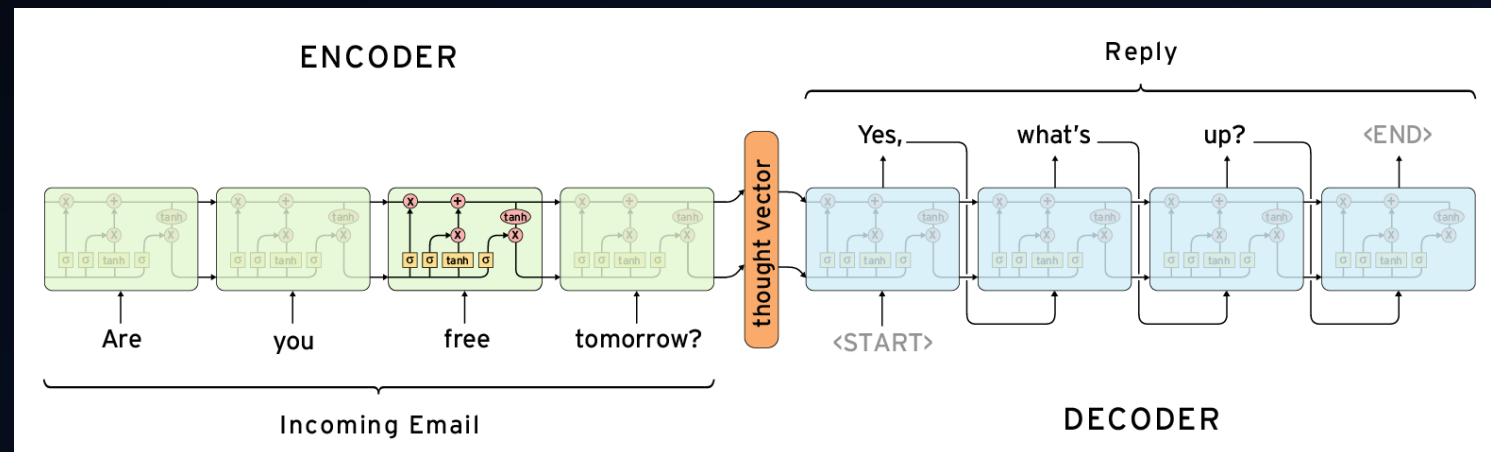
1. 2012: The Year of the Deep Net
2. 2013: The Year of Atari
3. 2016: The Year of AlphaGo
4. Deep Learning in Everyday Technology
5. Frontiers in Deep Learning



# Outline

## 3. Selected Modern Deep Learning Approaches

1. Residual Networks
2. Sequence to Sequence Models
3. Attention Mechanisms
4. GANs

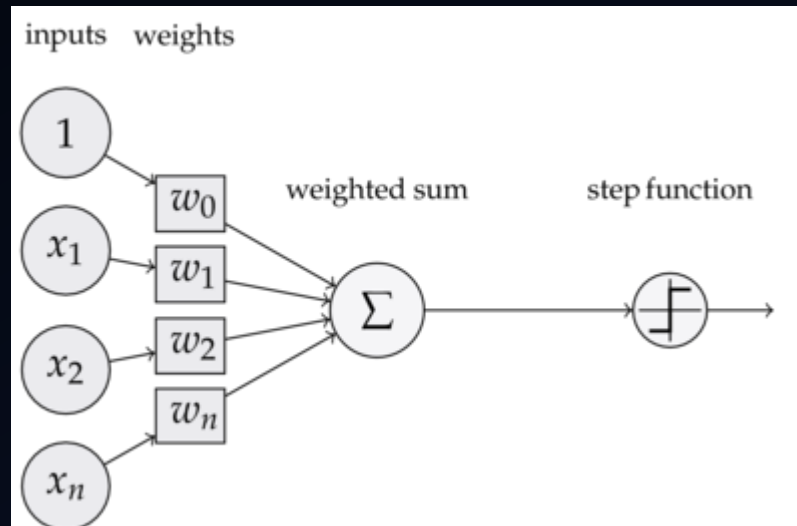





# Fundamentals of Neural Nets

# Early AI: the Perceptron


- Late 1950s
- Implemented on custom hardware
- Linear classification
- Awkward training scheme based on error



$$y = h(w^T x + b) = \begin{cases} 1 & w^T x + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

# The Perceptron


- Interpret as checking for presence of a pattern and issuing binary response
- Inspired by neurons in brain
- **Can only learn very simple functions: not even expressive enough to learn XOR**





# Early Neural Networks

- Neural networks emerged as differentiable, **multi-layered** perceptrons





# Neural Networks

- A model composed of several “layers,” where the ones between input and output are called “hidden”
- Standard basic layer:
  - Linearity,  $z = Wx + b$
  - Followed by elementwise nonlinearity,  $y = f(z)$
- With  $L$  layers:

$$\begin{aligned} a^0 &= x \\ a^j &= f^j(W^j a^{j-1} + b^j) \\ y &= a^L \end{aligned}$$


# Neural Networks

- Common nonlinearities (activation functions):

- Sigmoid  $\sigma(x) = \frac{1}{1+e^{-x}}$

- Tanh  $\tanh(x) = \frac{1-e^{-2x}}{1+e^{-2x}}$

- Relu  $\text{relu}(x) = \max(0, x)$



# Neural Networks

- Common output layers:
  - No activation (good for regression)
  - Softmax activation (good for classification): with  $z$  the output of the last linear transform,

$$y_i = \frac{\exp(z_i)}{\sum_{j=1}^K \exp(z_j)}$$

This gives a **probability distribution of dimension K** as an output

# Neural Networks

- A neural network is characterized by its **architecture** and its **parameters**
- Architecture (fixed)
  - A term used to describe design choices, like
    - Number of layers
    - Size of output from each layer
    - Nonlinearities
    - Alternative connection patterns
- Parameters (learned)
  - “weights”  $W$  and “biases”  $b$  at each layer, and possibly others

# Neural Networks

- What do you do with them?
  - Classification
  - Prediction / forecasting
  - Function approximation
  - Decision-making of any kind, really
- How do you train them (learn the parameters)?
  - Most powerful tool of all: **gradient descent**

# Training Neural Networks: Task and Loss Function

- Suppose that  $f_\theta$  is your neural net, and  $\theta = (W^1, \dots, W^L, b^1, \dots, b^L)$  is your set of parameters
- Task is classification
- You have a dataset of pairs  $D = \{(x^i, y^i)\}_{i=1, \dots, N}$

- Form a per-datum loss function:

$$L(\theta, x, y) = d(f_\theta(x), y)$$

- Form a **dataset average loss function**:

$$\mathcal{L}(\theta, D) = \frac{1}{N} \sum_{i=1}^N L(\theta, x^i, y^i)$$

# Training Neural Networks: Loss Function

- Suppose  $y \in \{1, \dots, K\}$ , and the last layer of  $f_\theta$  is softmax
- Choose differentiable per-datum loss function, like **cross-entropy loss**:

$$L(\theta, x, y) = - \sum_{i=1}^K \mathbb{I}[y = i] \log P(y = i | x, \theta)$$

$$\text{where } P(y = i | x, \theta) = [f_\theta(x)]_i$$

- Dataset average loss is now differentiable with respect to  $\theta$



# Training Neural Networks: Gradient Descent and SGD

- Gradient descent:

$$g = \nabla_{\theta} \mathcal{L}(\theta, D)$$
$$\theta \leftarrow \theta - \alpha g$$

- Problem: expensive to compute gradient of loss over whole dataset
- Solution: **stochastic gradient descent (SGD)**

$$g \approx \sum_{(x,y) \in B} L(\theta, x, y)$$

where  $B$  is randomly-sampled **minibatch**

# Training Neural Networks: Backpropagation

- In the 1970s people rediscovered the wheel and gave it a new name
- Algorithm for computing gradients of neural network outputs with respect to parameters is called **backprop**

# Training Neural Networks: Backpropagation

Chain rule:

$$\frac{\partial y}{\partial W^j} = \frac{\partial y}{\partial a^L} \frac{\partial a^L}{\partial a^{L-1}} \cdots \frac{\partial a^{j+1}}{\partial a^j} \frac{\partial a^j}{\partial W^j}$$

Define

$$\delta^j = \frac{\partial y}{\partial a^L} \frac{\partial a^L}{\partial a^{L-1}} \cdots \frac{\partial a^{j+1}}{\partial a^j}$$

If you have  $\delta^L$ , easy to perform “backprop”:

For  $j = L - 1, L - 2, \dots, 1$ :

$$\delta^j = \delta^{j+1} \frac{\partial a^{j+1}}{\partial a^j}$$

$$\frac{\partial y}{\partial W^j} = \delta^j \frac{\partial a^j}{\partial W^j}$$

# Training Neural Networks: SGD Variants

- SGD stops updating when gradient is small---this should happen at loss function minima, but also happens in **valleys**

- Momentum methods break out of valleys

$$\begin{aligned}v &\leftarrow \gamma v + g \\ \theta &\leftarrow \theta - \alpha v\end{aligned}$$

- Adaptive learning rate algorithms can also help, like RMSprop:

$$\begin{aligned}v &\leftarrow \gamma v + (1 - \gamma)g^2 \\ \theta &\leftarrow \theta - \alpha \frac{g}{\sqrt{v} + \epsilon}\end{aligned}$$

# Special Neural Net Architectures: Convolutional Layers

- Layers we described previously are *fully-connected*
- That is bad for parameter reuse
- How can we encode **invariance** to translation?



# Special Neural Net Architectures: Convolutional Layers

- Conv layers *convolve* an input with some feature kernel to produce a *response map*

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

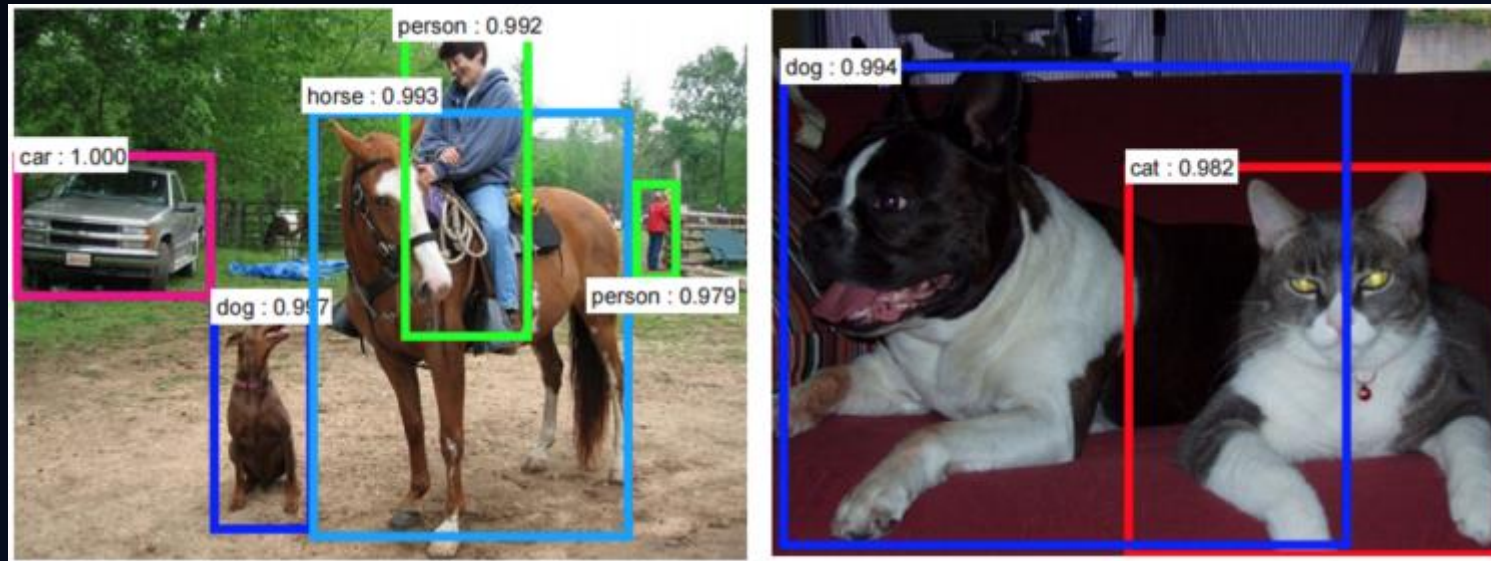
Image

4		

Convolved  
Feature

# Special Neural Net Architectures: Convolutional Layers


- Conv layers are great for parameter reuse
- Far fewer parameters than fully-connected layers
- Excellent at vision---can learn to understand contents of images from raw pixels with no other feature extraction!





# Special Neural Net Architectures: Recurrent Nets

- Some tasks require **memory** to solve
  - Time series prediction
  - Language modeling
- Recurrent Neural Networks (RNNs) **have memory in hidden state**



# Special Neural Net Architectures: Recurrent Nets

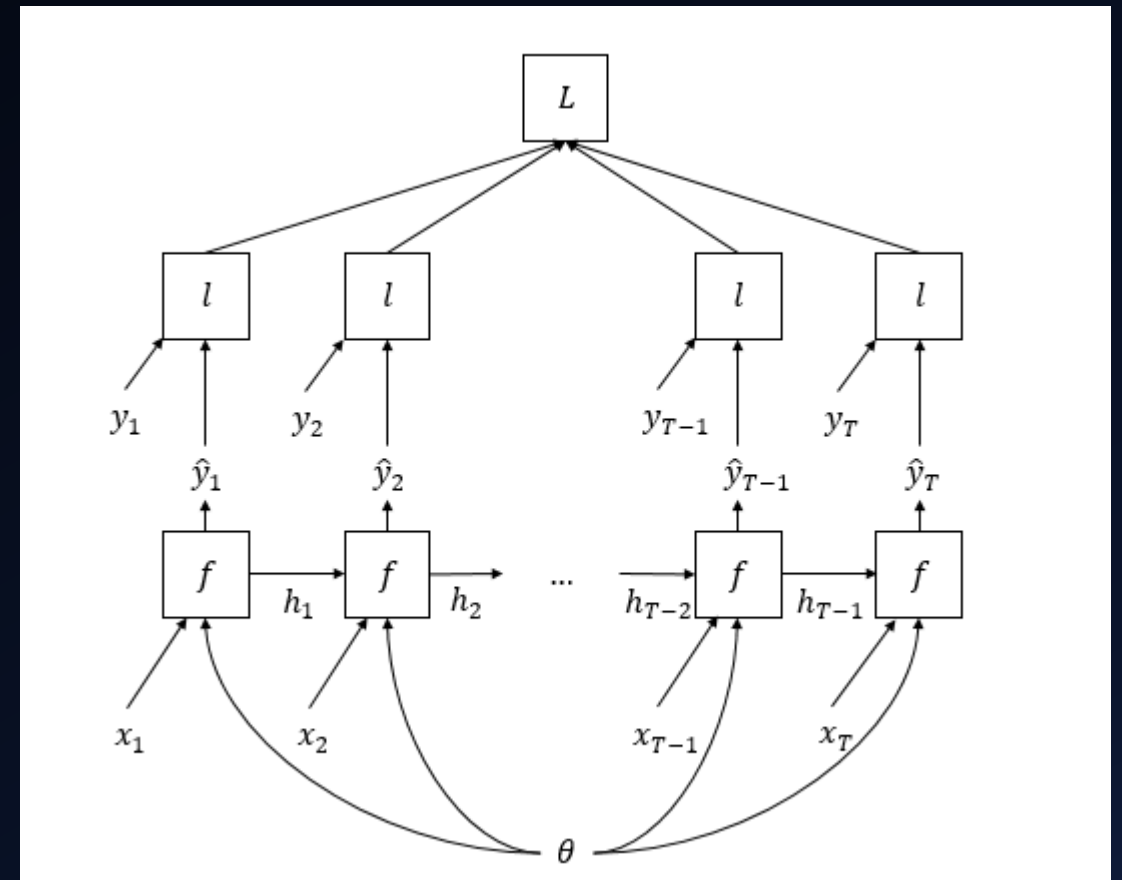
- “Vanilla” RNN layer:

$$h_t = \sigma(Wx_t + Rh_{t-1} + b)$$

- Drop in as a replacement to a fully-connected layer
- Can stack RNN layers in the same way as standard layers
- Requires special training: **backprop through time (BPTT)**

# Special Neural Net Architectures: Recurrent Nets

- Backprop Through Time:
  - Loss function on **sequences**: per-sequence loss is a sum or average of per-time step losses
  - $(x_1, \dots, x_T, y_1, \dots, y_T)$
  - Need dataset of sequences
  - If sequences are too long, **truncate** backprop to some horizon



# Vanishing Gradients in Vanilla RNNs

- Ability to learn long time dependencies requires **good gradient flow through network**
- Problem:

$$\frac{\partial h_t}{\partial h_k} = \frac{\partial h_t}{\partial h_{t-1}} \frac{\partial h_{t-1}}{\partial h_{t-2}} \cdots \frac{\partial h_{k+1}}{\partial h_k}$$

$$\frac{\partial h_t}{\partial h_{t-1}} \propto R$$

$$\frac{\partial h_t}{\partial h_k} \propto R^{t-k}$$

- $\lambda(R)$  large  $\rightarrow$  Gradient explodes!
- $\lambda(R)$  small  $\rightarrow$  Gradient vanishes...


# Solution to Vanishing Gradients: LSTMs

- Long Short-Term Memory (LSTM) Networks
  - Have two hidden states
  - One of which is *additive* instead of multiplicative, making gradient flow easier:

$$c_t = c_{t-1} + z_t$$

(in truth, also have forget gates and input gates – omitted here for simplicity)


- Allow learning of dependencies over **hundreds of time steps!**



# Recent History of Deep Learning

# 2012 – The Year of the Deep Net

- After years of many incremental advances in neural nets (development of ReLU and advanced training techniques), a deep neural network trained end-to-end by SGD won ImageNet contest
- AlexNet (Krizhevsky et al.):





# 2012 – The Year of the Deep Net

- ImageNet Large Scale Visual Recognition Contest (ILSVRC)
  - ~1000 high-resolution images in each of 1000 categories
- AlexNet substantially improved SOTA using no specialized algorithms:

	Top-1 Error Rate	Top-5 Error Rate
Sparse Coding	47.1	28.2
SIFT+FVs	45.7	25.7
<b>AlexNet</b>	<b>37.5</b>	<b>17.0</b>

- Interest in deep convolutional networks grew substantially

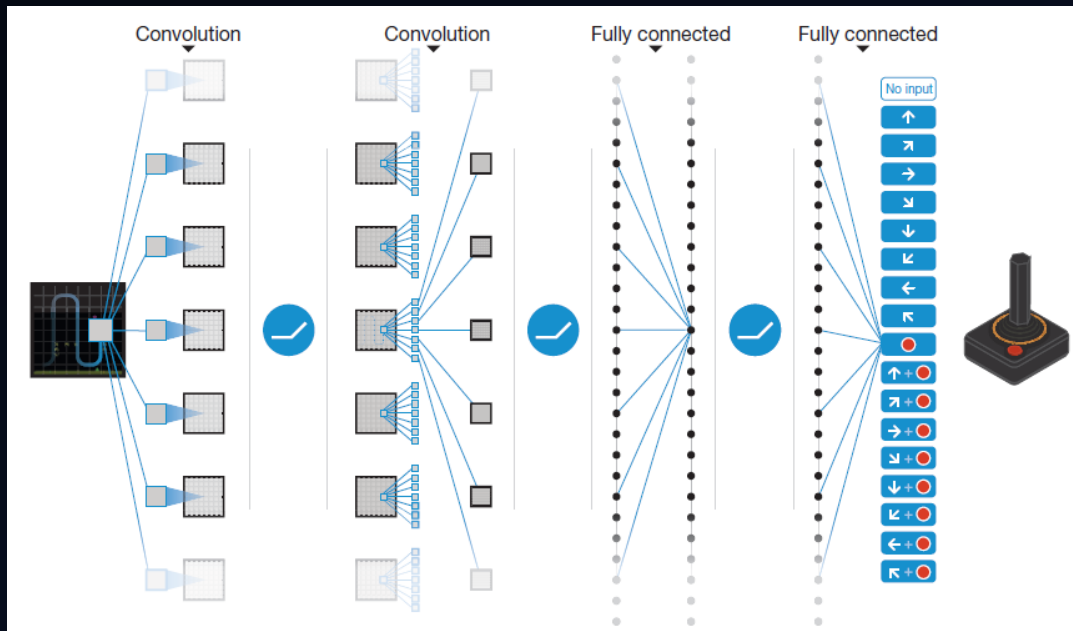
# 2012 – The Year of the Deep Net

- Le et al. (folks at Google and Stanford) trained a deep 9-layer **autoencoder** on a huge dataset from YouTube
- Autoencoder loss:  $d(x, f_{\theta}(x))$
- *Unsupervised* learning  
(just need data, no labels)
- Learned neurons that detected human faces, bodies, cat faces, etc.



# 2013 – The Year of Atari

- A convolutional neural network was trained to play Atari games by **reinforcement learning (RL)**
- Mnih et al. introduced **Deep Q-Learning**, first major breakthrough in RL + deep learning



# Deep Q-Learning

- Goal is to learn **optimal action-value function**  $Q^*(s, a)$
- Bellman equation (dynamic programming!):


$$Q^*(s, a) = \mathbb{E}_{s'}[r(s, a) + \gamma \max_{a'} Q^*(s', a')]$$

- Learn approximator  $Q_\theta$  by minimizing **squared Bellman error**

$$L(\theta, D) = \mathbb{E}_{s,a,r,s' \sim D} \left[ \left( Q_\theta(s, a) - \left( r + \gamma \max_{a'} Q_\theta(s', a') \right) \right)^2 \right]$$

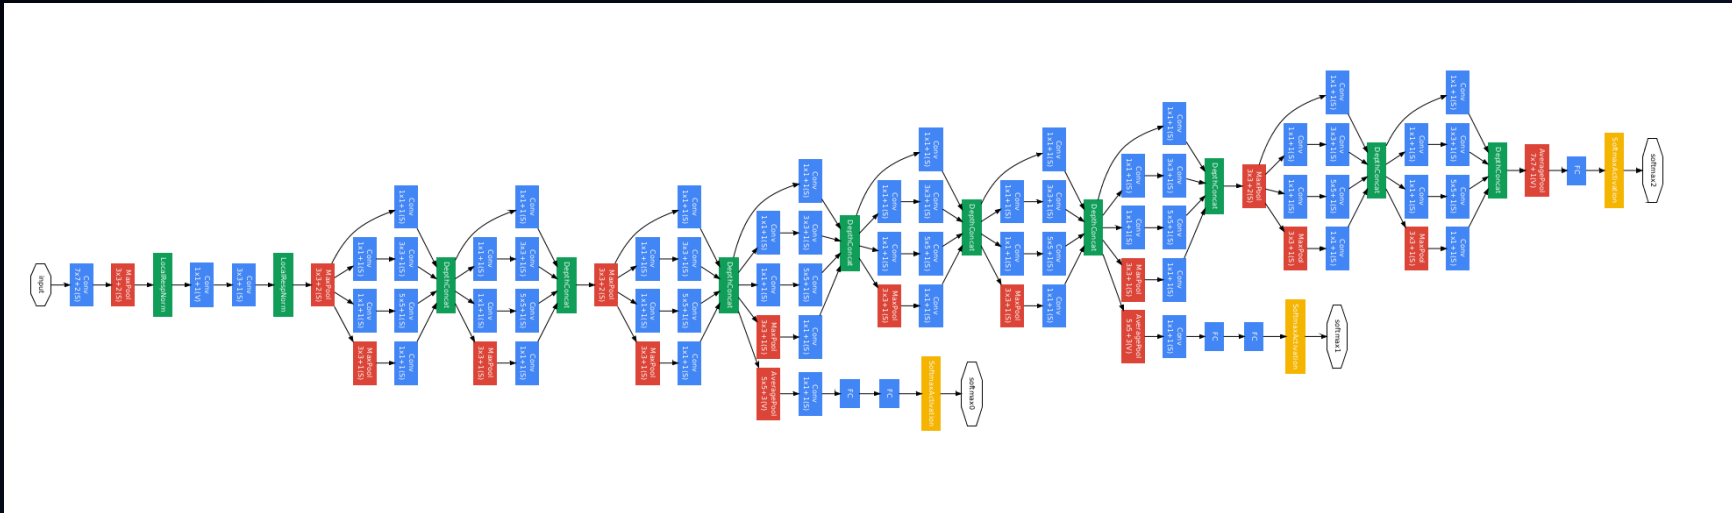
# Deep Q-Learning

- Achieved **superhuman performance** on several Atari games!
- Many improvements followed
- Shown right: results from 2015 Nature paper on DQN



# 2014 – GoogLeNet

- GoogLeNet (Szegedy et al.): 22-layer deep convolutional network gets **human-level performance on ImageNet**



- Top-5 error rate of **6.67%**
- Every blue block is a conv layer!

# 2016 – The Year of AlphaGo


- Go is a game with simple rules and extremely complex play
- Impossible to solve with minimax AI / hard-coded rules





# 2016 – The Year of AlphaGo

- Google DeepMind developed **AlphaGo**, a deep learning AI for Go
- Used 13-layer networks to select optimal moves and evaluate positions
- **Monte Carlo Tree Search** to help AlphaGo think ahead



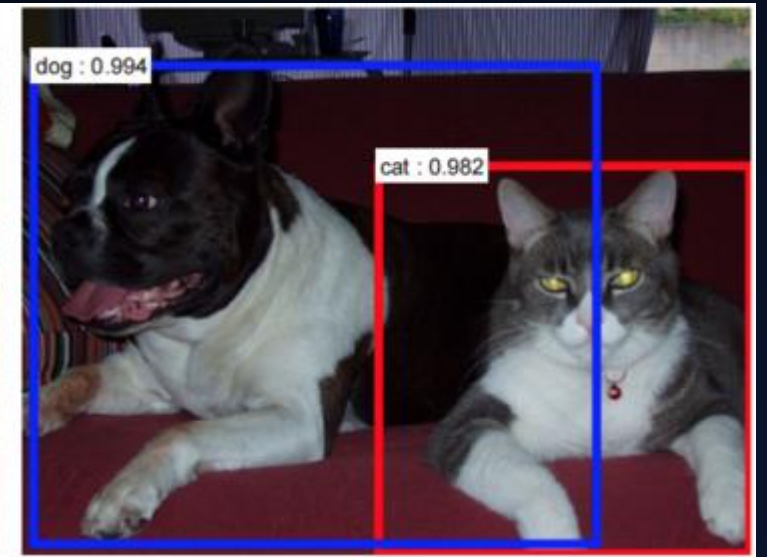
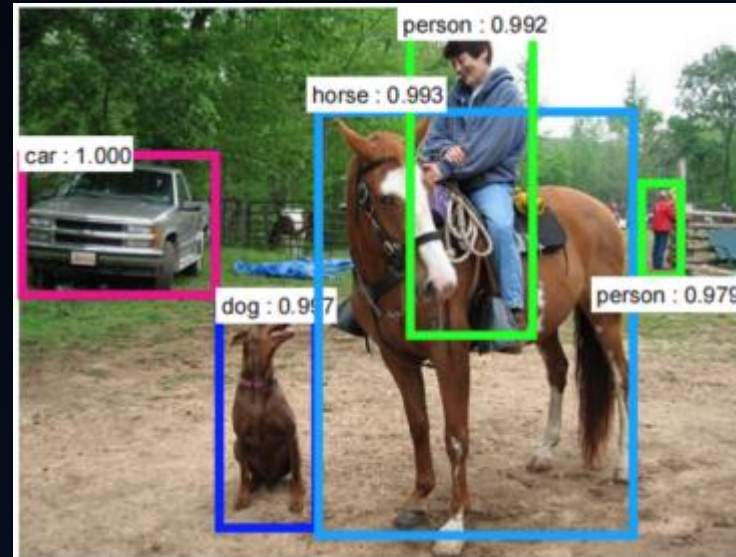

# 2016 – The Year of AlphaGo

- Defeated human grand master, Lee Sedol, 4-1!



# Deep Learning in Everyday Technology


- Face recognition and object detection: convolutional nets



- Near future: vision for self-driving cars!

# Deep Learning in Everyday Technology

- Neural machine translation




<https://devblogs.nvidia.com/parallelforall/introduction-neural-machine-translation-gpus-part-3/>



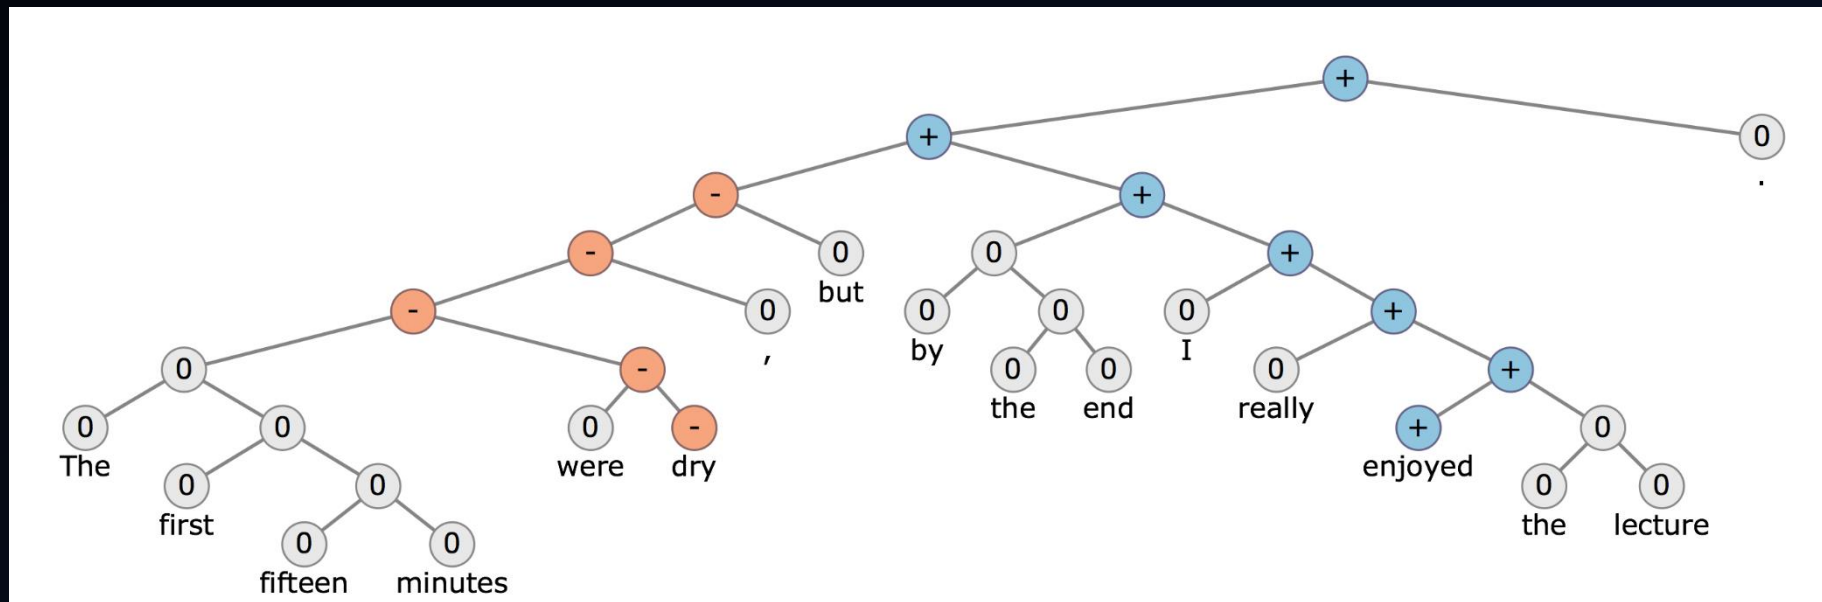
# Deep Learning in Everyday Technology

- Neural style transfer (Gatys et al.)



# Deep Learning in Everyday Technology

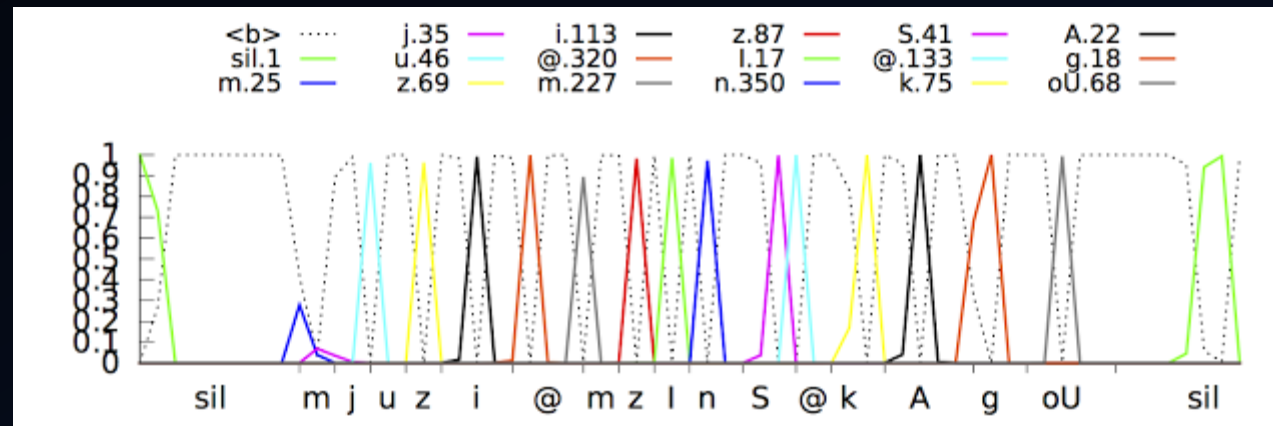
- Data analytics: sentiment analysis



(image from Stanford course on deep learning for NLP)

# Deep Learning in Everyday Technology

- Automatic speech recognition
  - Core technology for Google's speech recognition is based on deep learning!




<https://research.googleblog.com/2015/09/google-voice-search-faster-and-more.html>



# Frontiers in Deep Learning

- Speech synthesis (WaveNet)
- Robotics
- One-shot learning / Metalearning
  - Learning to Learn by Gradient Descent
  - RL<sup>2</sup>
  - MAML
- Chatbots





# Frontiers in Deep Learning: Chatbots

**Human:** *what is the purpose of life ?*

**Machine:** *to serve the greater good .*

**Human:** *what is the purpose of living ?*

**Machine:** *to live forever .*

**Human:** *what is the purpose of existence ?*

**Machine:** *to find out what happens when we get to the planet earth .*

**Human:** *what is the purpose of emotions ?*

**Machine:** *i don 't know .*

(Vinyals and Le, 2015)



# Selected Modern Deep Learning Approaches

# Residual Networks

- He et al., 2015, used **residual networks** to achieve 3.57% top-5 error on ImageNet LSVRC


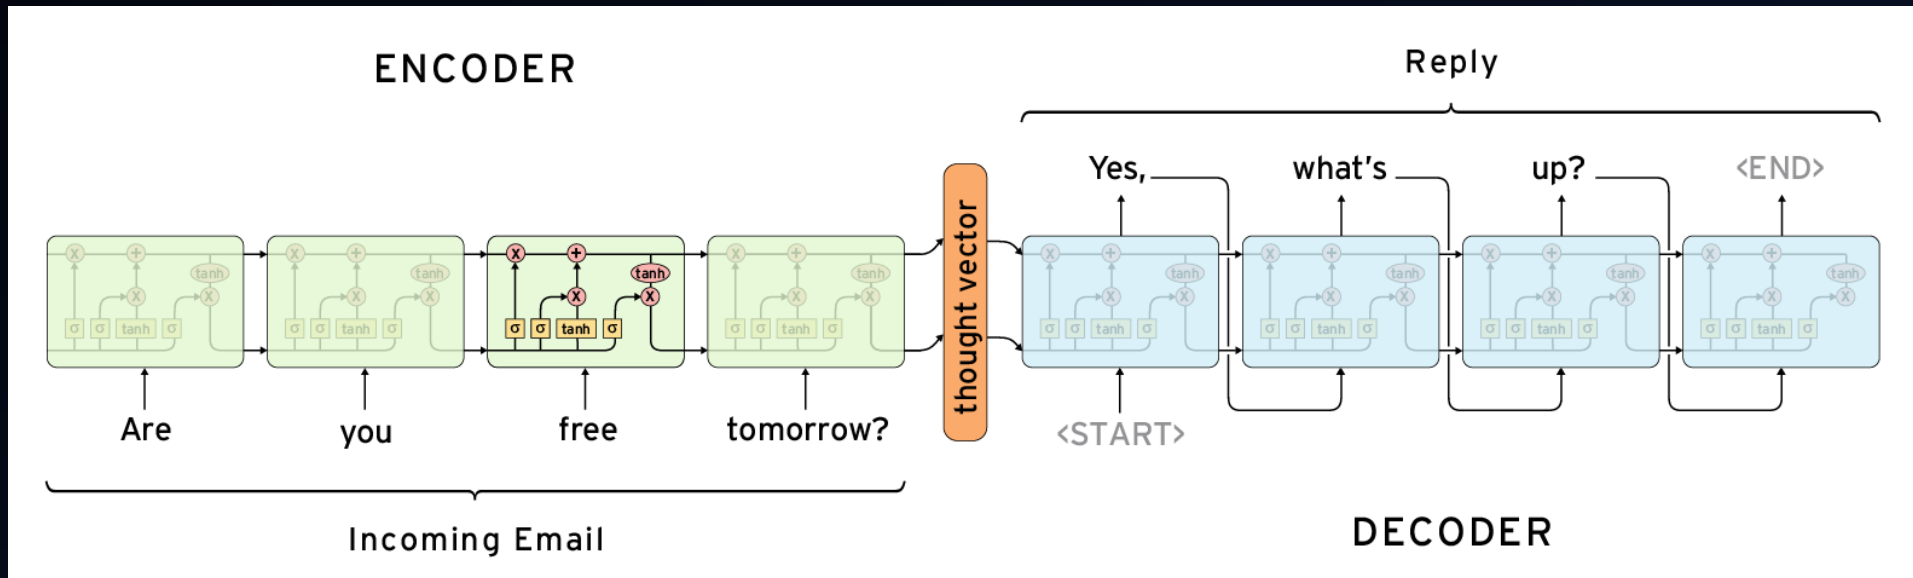


Figure 2. Residual learning: a building block.

# Sequence to Sequence Models

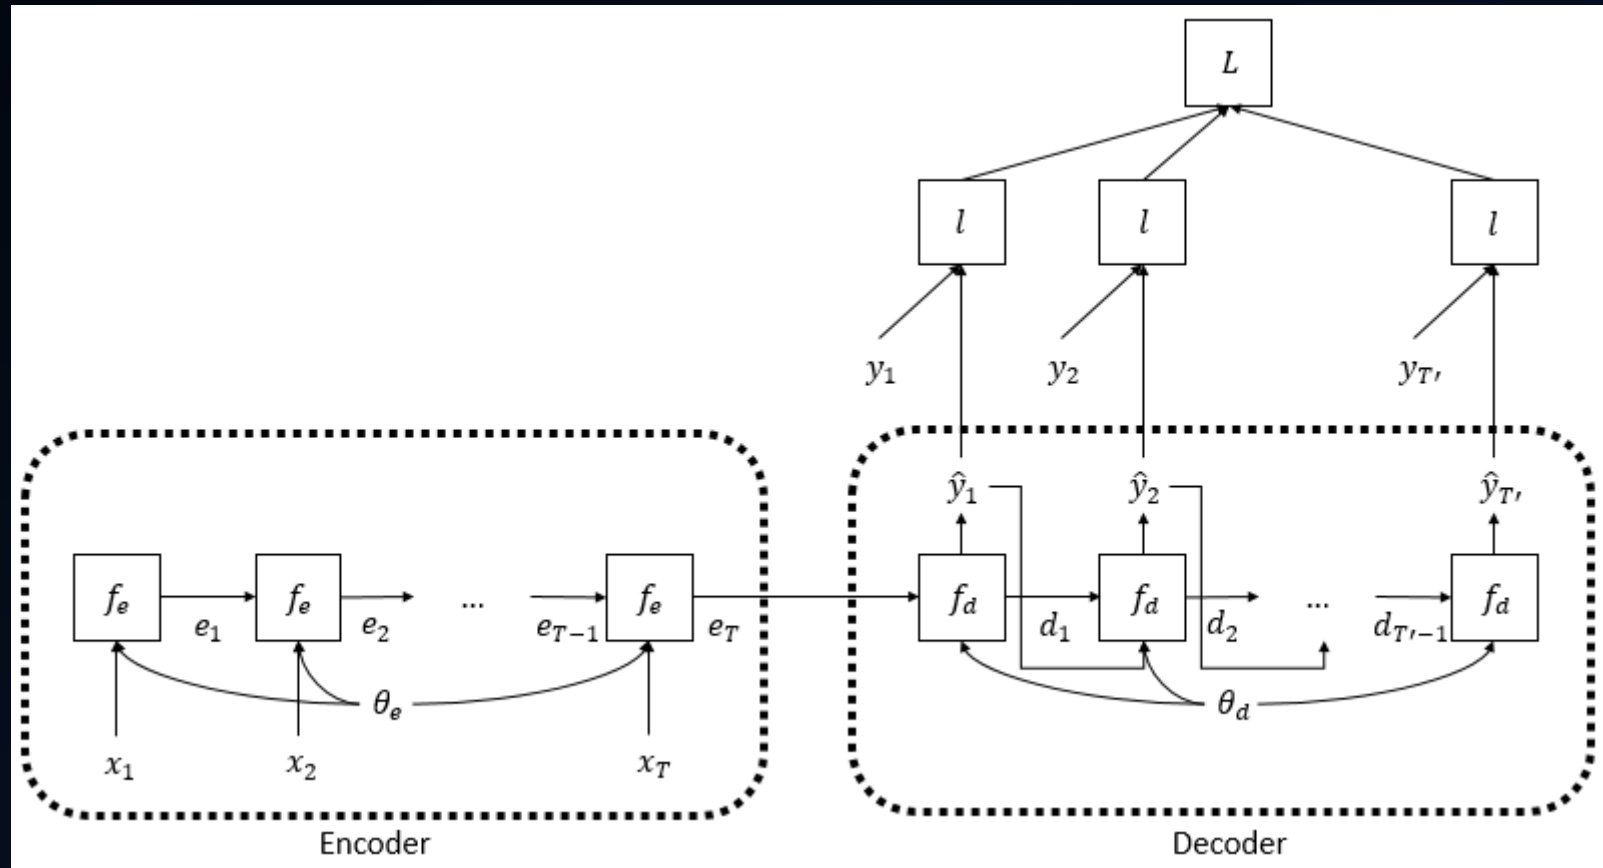
- Cho et al 2014, Sutskever et al 2014
- Key for neural machine translation



<http://suriyadeepan.github.io/2016-12-31-practical-seq2seq/>

# Sequence to Sequence Models

- How do you train them? Per-(input sequence, output sequence) pair loss function




# Attention Mechanisms

- Help sequence to sequence models by letting them focus on specific input (Bahdanau et al, 2015)
- For each encoder hidden state  $e_j$ , produce a weight  $w_j$ , and use

$$d' = \sum_j w_j e_j$$

as auxiliary input



# Generative Adversarial Networks (GANs)

- Goodfellow et al., 2014
- Goal is to learn a **generative model**
- IE, put noise in, get sample from data distribution out
- Learn in **adversarial game** with two players
  - Generator: trying to produce outputs like data
  - Discriminator: trying to tell if input is real or from generator

$$\min_G \max_D \mathbb{E}_{x \sim p_x} [\log D(x)] + \mathbb{E}_{z \sim p_z} [\log (1 - D(G(z)))]$$



# Generative Adversarial Networks (GANs)

- Can generate almost-natural looking images, but things get weird sometimes
- Hard to train, but this is an area of active research!





# Generative Adversarial Networks (GANs)

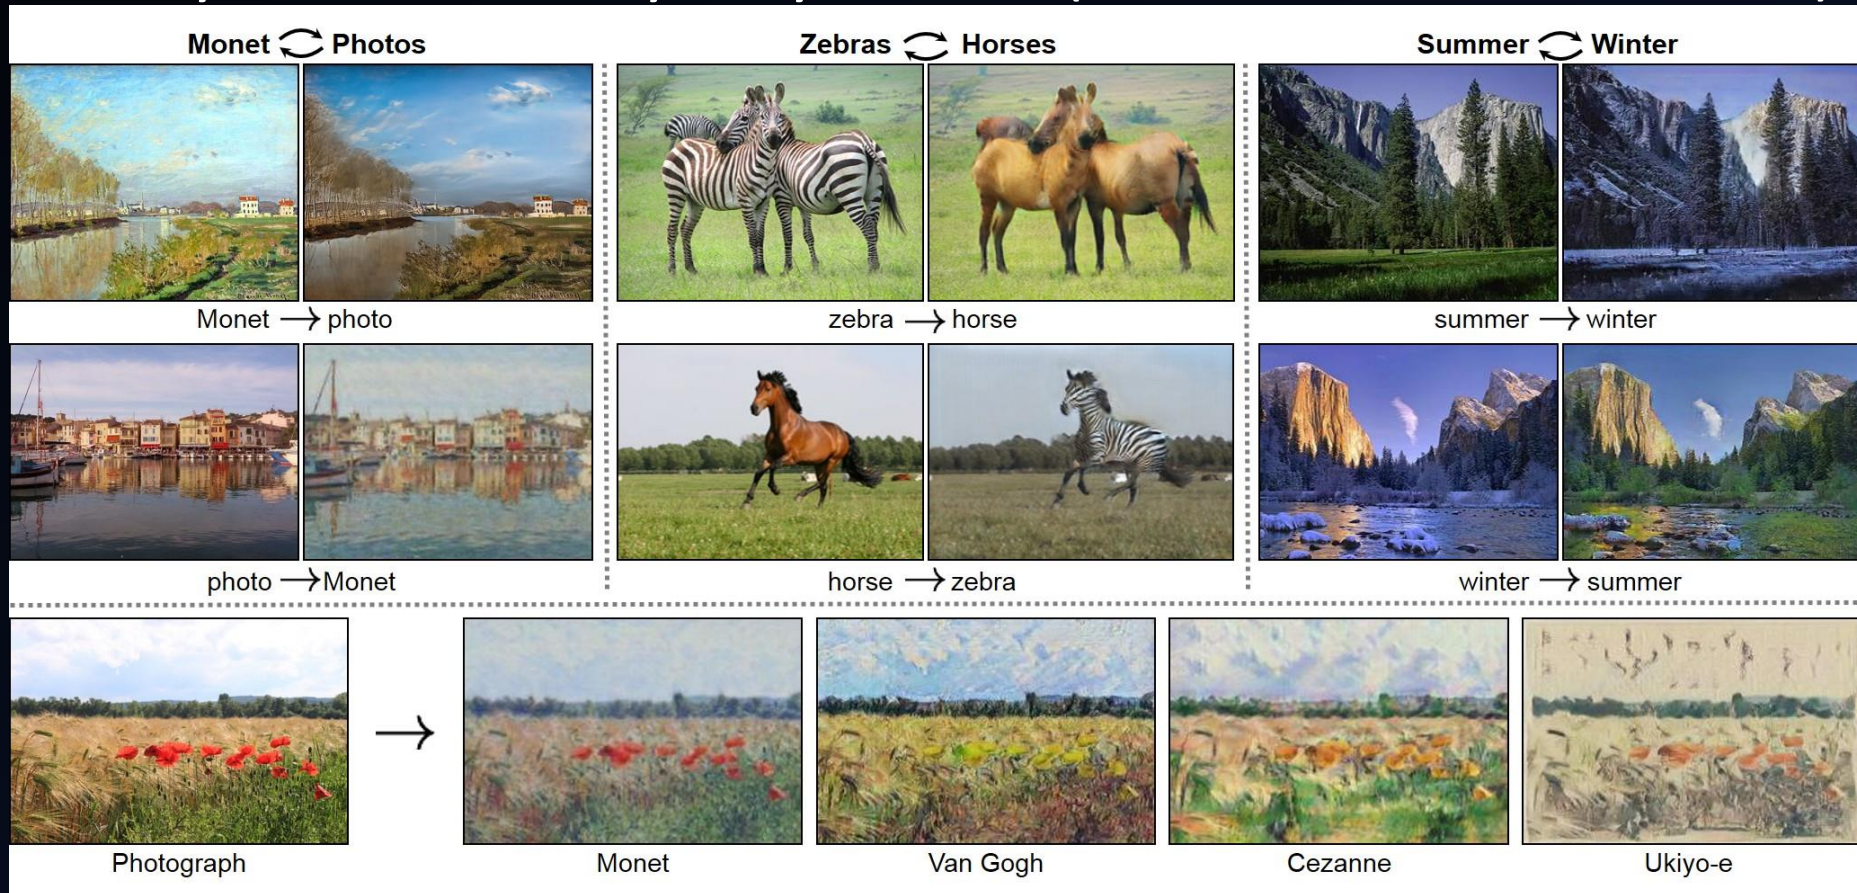
- Many variants, including generating images from **natural language text** (Reed et al. 2016)

this small bird has a pink  
breast and crown, and black  
primaries and secondaries.



# Generative Adversarial Networks (GANs)

- GAN + Cycle consistency = CycleGAN (Zhu and Park et al, 2017)



# And so much more!

- This talk only scratched the surface
- If you want, it's easy to get started learning...
  - Stanford Deep Learning Tutorial (<http://deeplearning.stanford.edu/>)
  - Andrej Karpathy Blog (<http://karpathy.github.io/>)
  - Deep Learning Textbook (<http://www.deeplearningbook.org/>)
  - Awesome Deep Learning Resource List (<https://github.com/ChristosChristofidis/awesome-deep-learning>)

# And so much more!

- ...and easy to get started hacking!
  - Tensorflow (<https://www.tensorflow.org/>)
  - Keras (<https://keras.io/>)
  - PyTorch (<http://pytorch.org/>)
  - Caffe (<https://caffe2.ai/>)
  - Chainer (<https://chainer.org/>)